

# Inside Hive (for beginners)

Takeshi NAKANO / Recruit Co. Ltd.

1

## Why?

- Hive is good tool for non-specialist!
- The number of M/R controls the Hive processing time.
  - ↓
  - How can we reduce the number?
  - What can we do for this on writing HiveQL?
    - ↓
    - How does Hive convert HiveQL to M/R jobs?
    - On this, what optimizing processes are adopted?

# Don't you have..

- This fb's paper has a lot of information!
- But this is a little old..

## Hive - A Warehousing Solution Over a Map-Reduce Framework

Ashish Thusoo, Joydeep Sen Sarma, Namit Jain, Zheng Shao,  
Prasad Chakka, Suresh Anthony, Hao Liu, Pete Wyckoff and Raghavam Murthy  
Facebook Data Infrastructure Team

### I. INTRODUCTION

The size of data sets being collected and analyzed in the industry for business intelligence is growing rapidly, making traditional warehousing solutions prohibitively expensive. *Hadoop* [3] is a popular open-source map-reduce implementation which is being used as an alternative to store and process extremely large data sets on commodity hardware. However, the map-reduce programming model is very low level and requires developers to write custom programs which are hard to maintain and reuse.

In this paper, we present *Hive*, an open-source data warehousing solution built on top of Hadoop. Hive supports queries expressed in a SQL-like declarative language - *HiveQL*.

databases. Each table has a corresponding HDFS directory. The data in a table is serialized and stored in files within that directory. Users can associate tables with the serialization format of the underlying data. Hive provides builtin serialization formats which exploit compression and lazy de-serialization. Users can also add support for new data formats by defining custom serialization and de-serialization methods (called *SerDe's*) written in Java. The serialization format of each table is stored in the system catalog and is automatically used by Hive during query compilation and execution. Hive also supports external tables on data stored in HDFS, NFS or local directories.

7/12/2011

HIVE - A warehouse solution over Map Reduce  
Framework

3

- Component Level Analysis

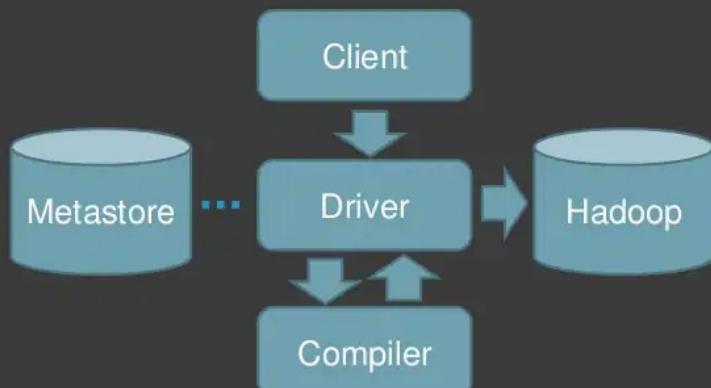


7/12/2011

HIVE - A warehouse solution over Map Reduce  
Framework

4

# Hive Architecture / Exec Flow



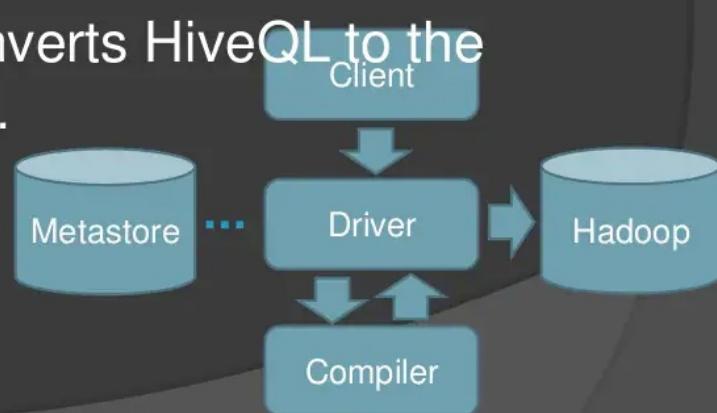
7/12/2011

HIVE - A warehouse solution over Map Reduce Framework

5

## Hive Workflow

- Hive has the operators which are minimum processing units.
- The process of each operator is done with HDFS operation or M/R jobs.
- The compiler converts HiveQL to the sets of operators.



7/12/2011

HIVE - A warehouse solution over Map Reduce Framework

6

# Hive Workflow

## ⦿ Operators

Operators	Descriptions
TableScanOperator	Read the data from table
ReduceSinkOperator	Build the result data to Reducer
JoinOperator	Join 2 data
SelectOperator	Reduce the output columns.
FileSinkOperator	Build the result data to output (file).
FilterOperator	Filter the input data.
GroupByOperator	...
MapJoinOperator	...
LimitOperator	...
UnionOperator	...

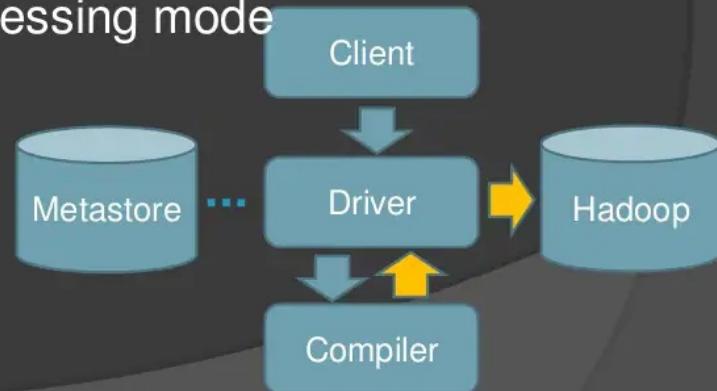
7/12/2011

HIVE - A warehouse solution over Map Reduce  
Framework

7

# Hive Workflow

- ⦿ For M/R processing, Hive uses ExecMapper and ExecReducer.
- ⦿ On processing, we have 2 modes.
  1. Local processing mode
  2. Distributed processing mode



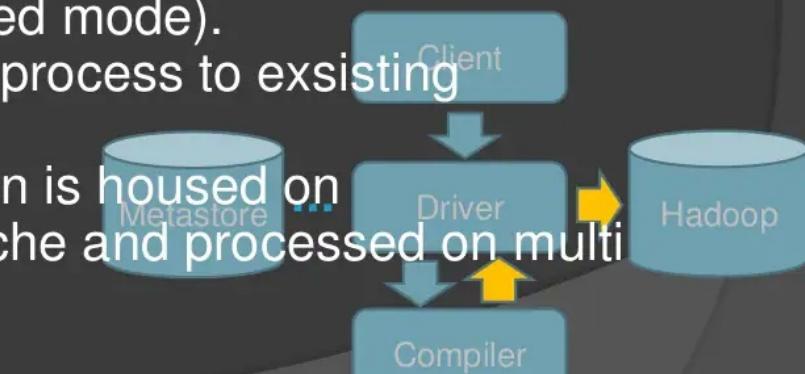
7/12/2011

HIVE - A warehouse solution over Map Reduce  
Framework

8

# Hive Workflow

- On 1(Local mode)  
Hive fork the process with hadoop command.  
The plan.xml is made just on 1 and the single node processes this.
- On 2(Distributed mode).  
Hive send the process to existing JobTracker.  
The information is housed on DistributedCache and processed on multi nodes.

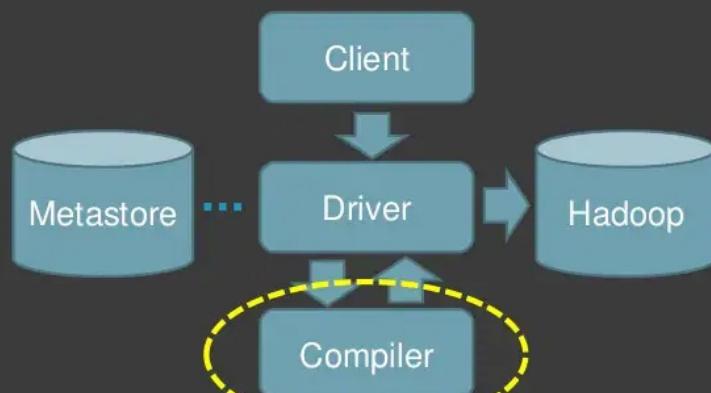


7/12/2011

HIVE - A warehouse solution over Map Reduce Framework

9

- Compiler : How to Process HiveQL

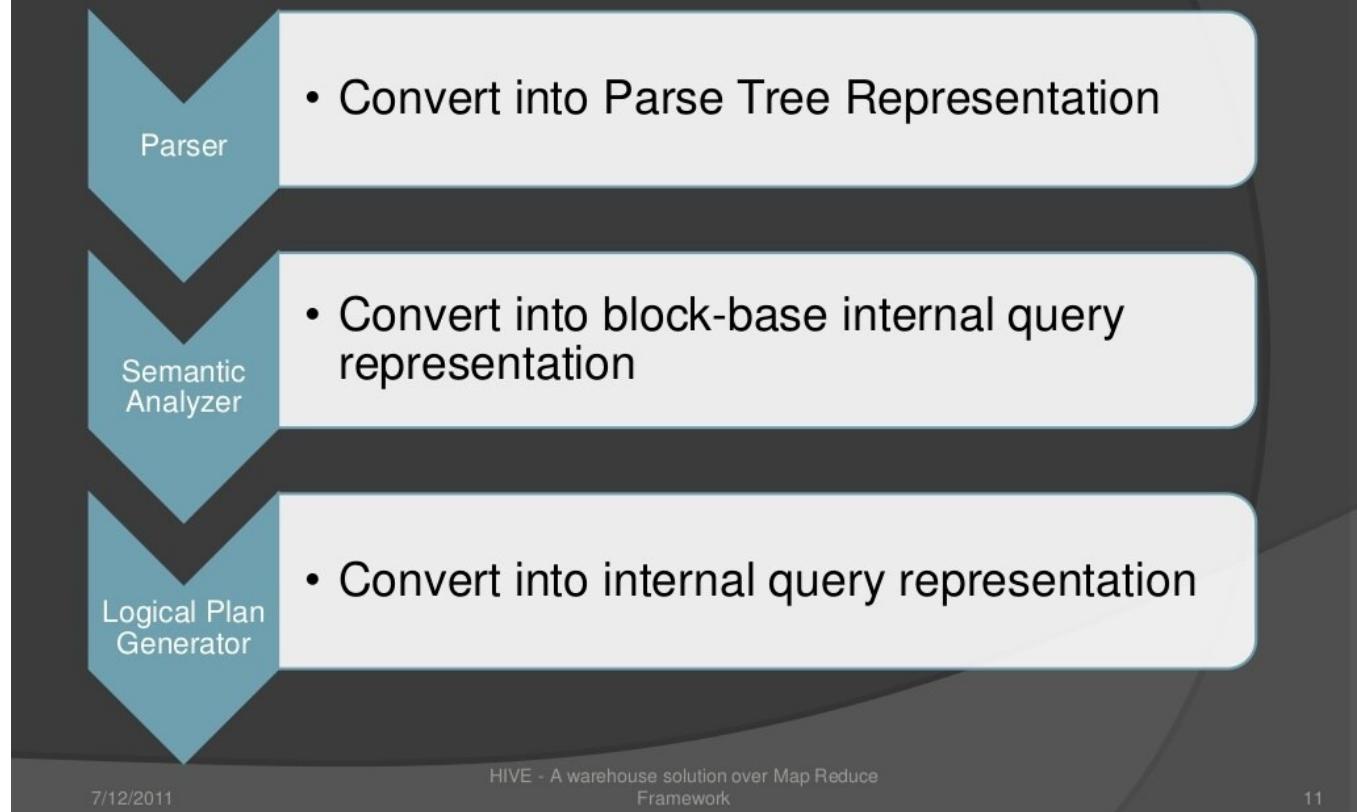


7/12/2011

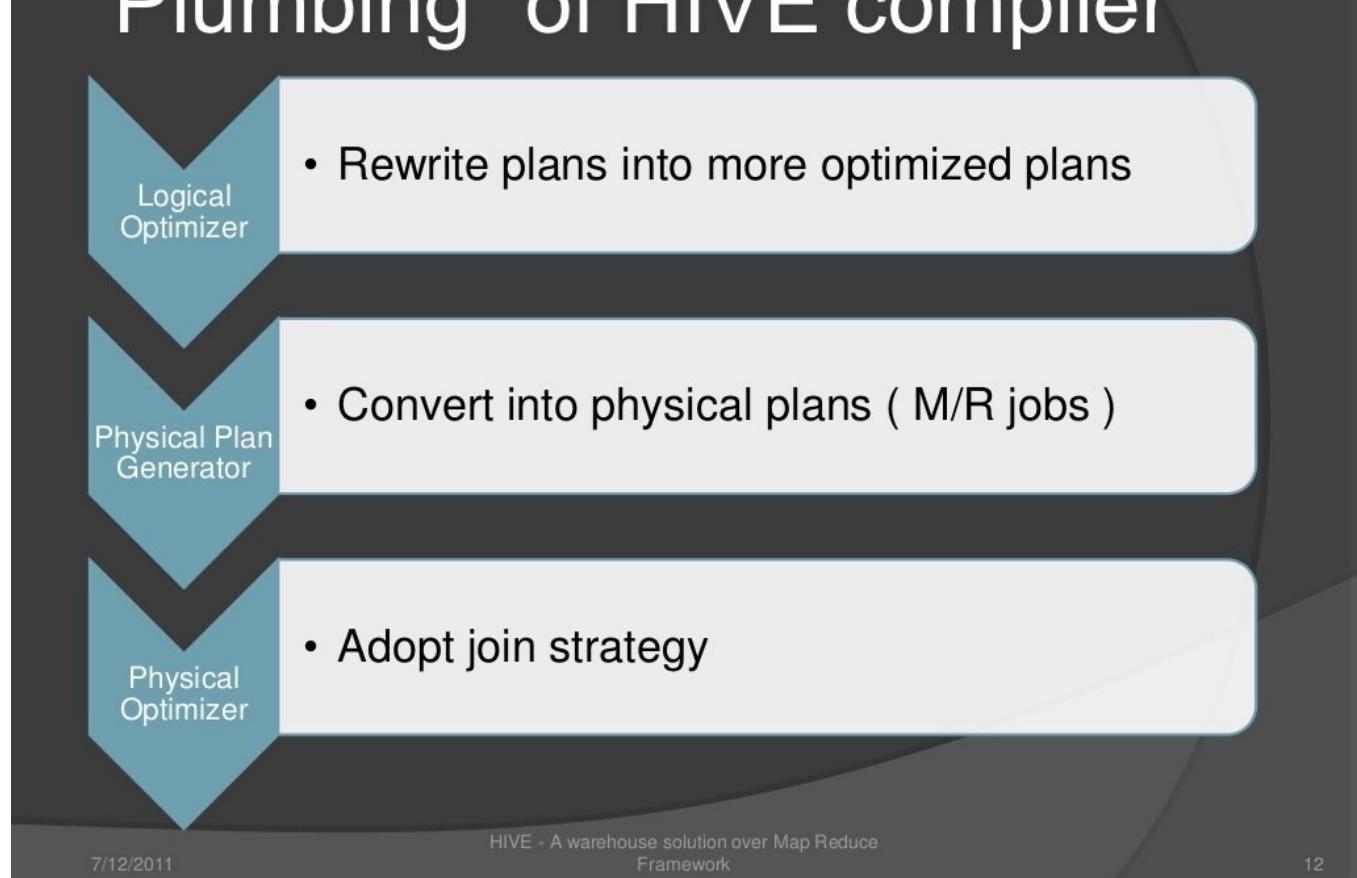
HIVE - A warehouse solution over Map Reduce Framework

10

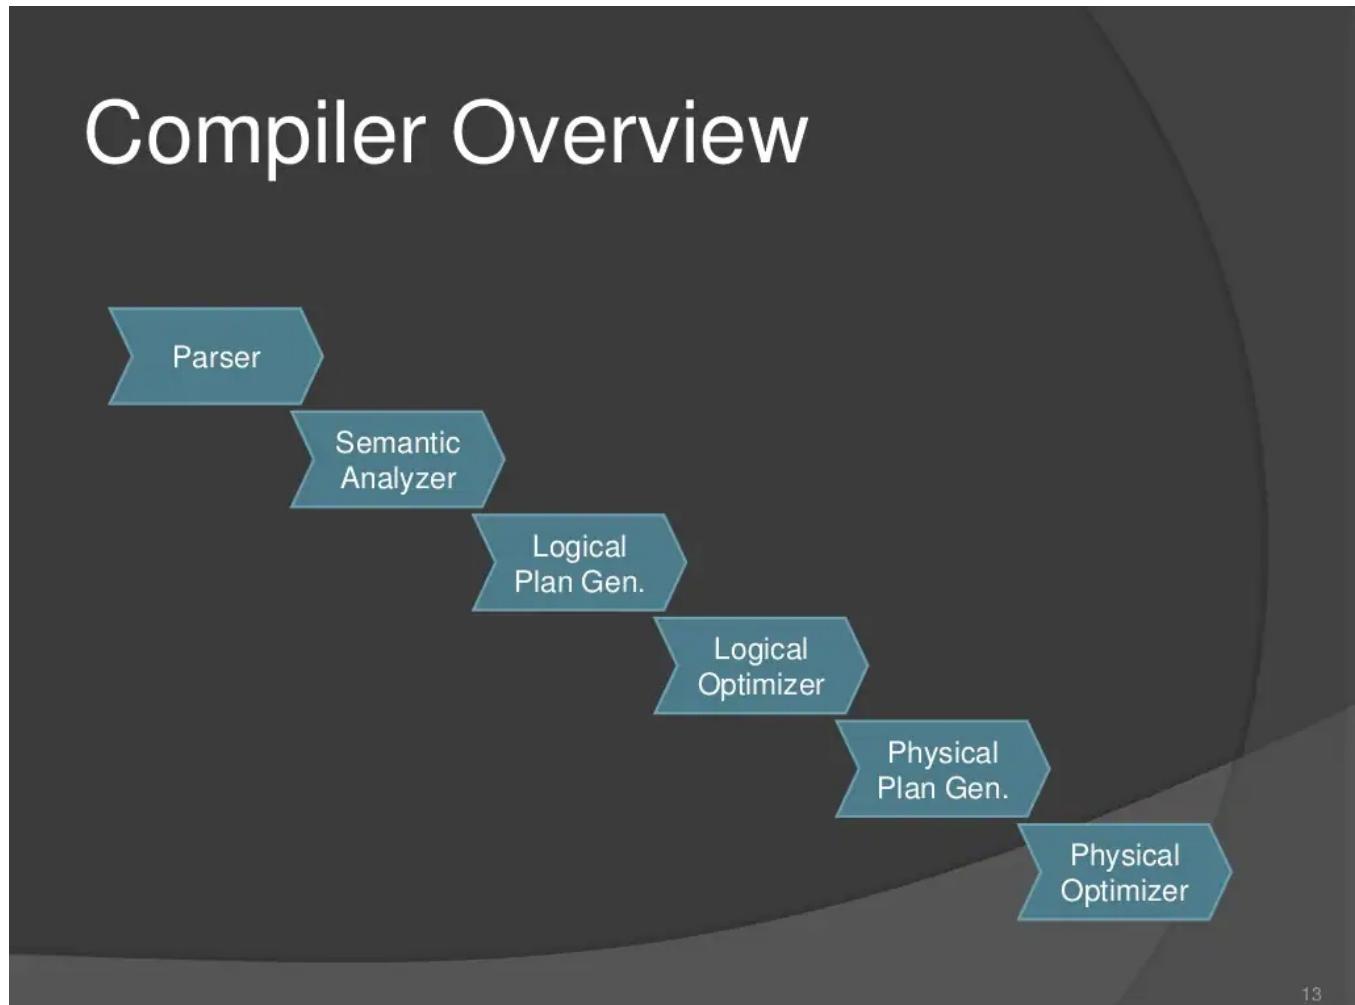
# “Plumbing” of HIVE compiler



# “Plumbing” of HIVE compiler

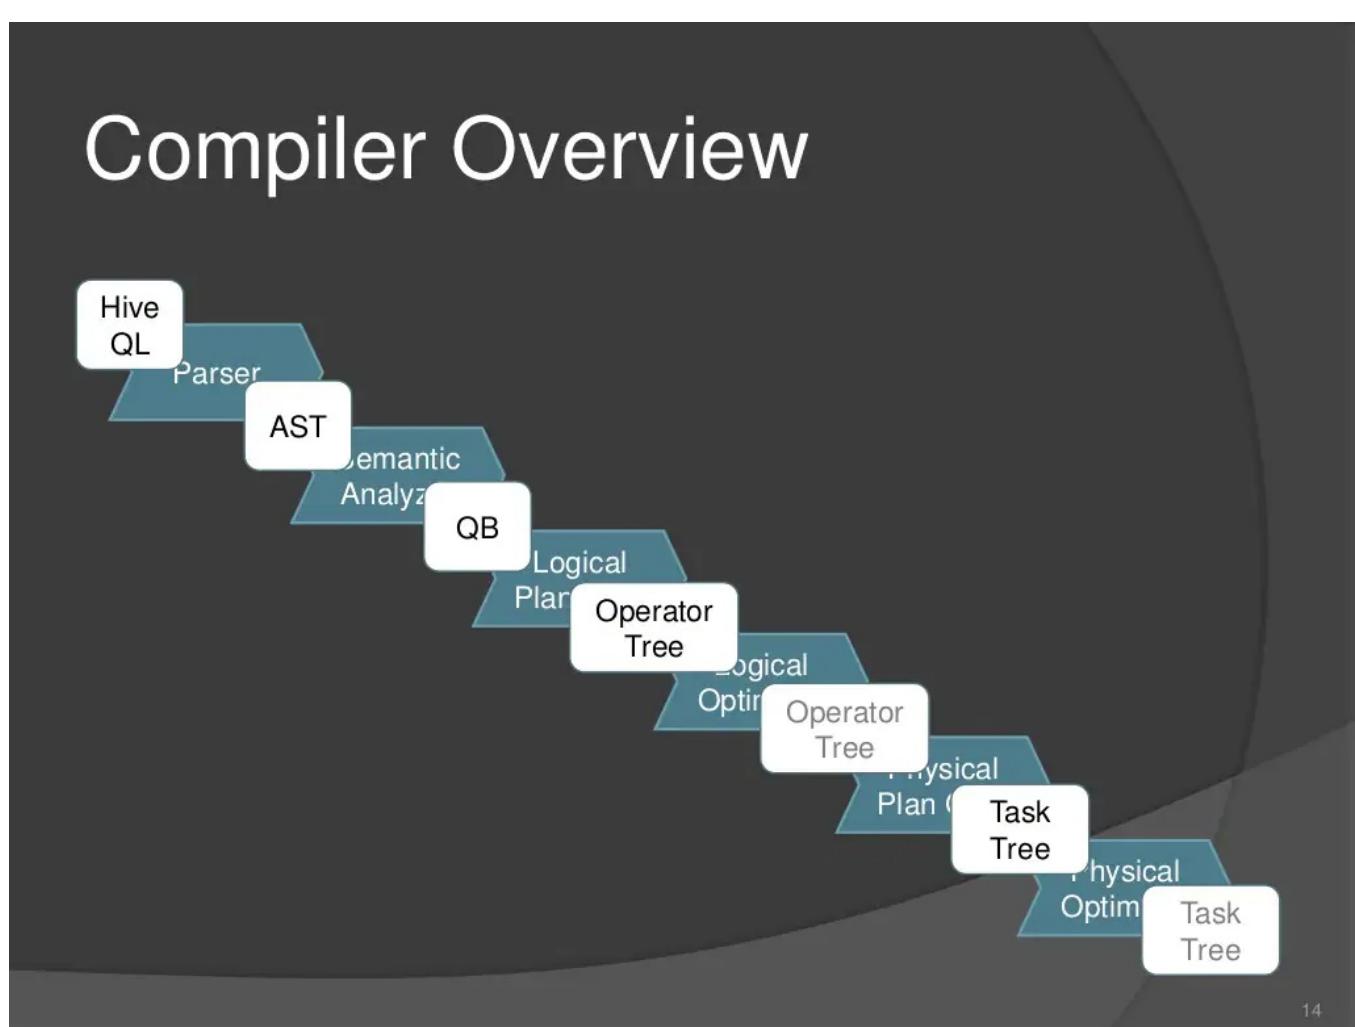


# Compiler Overview

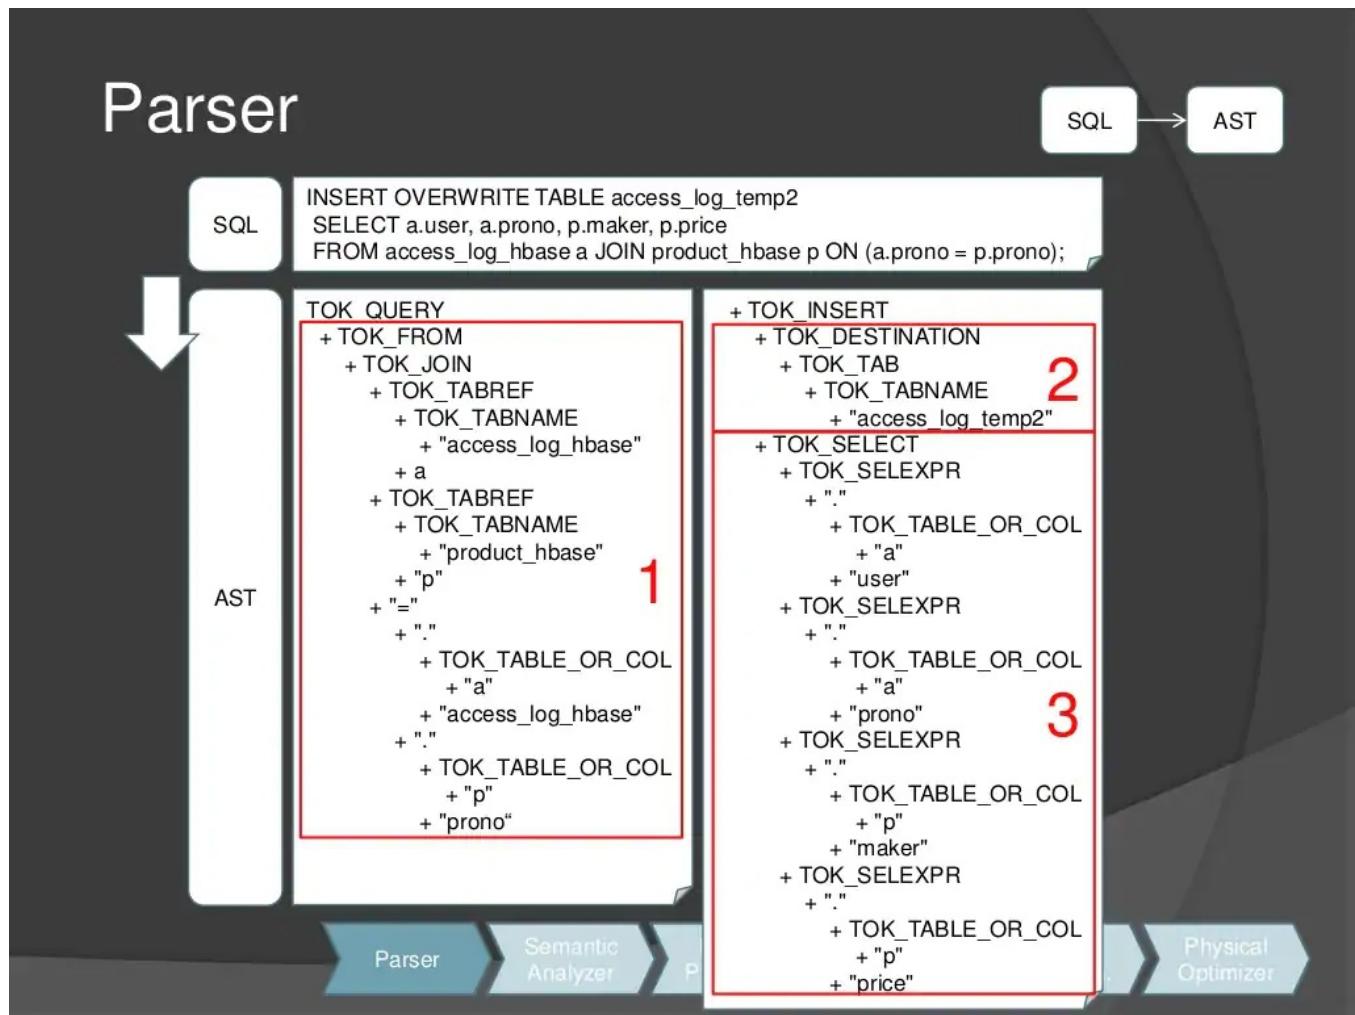
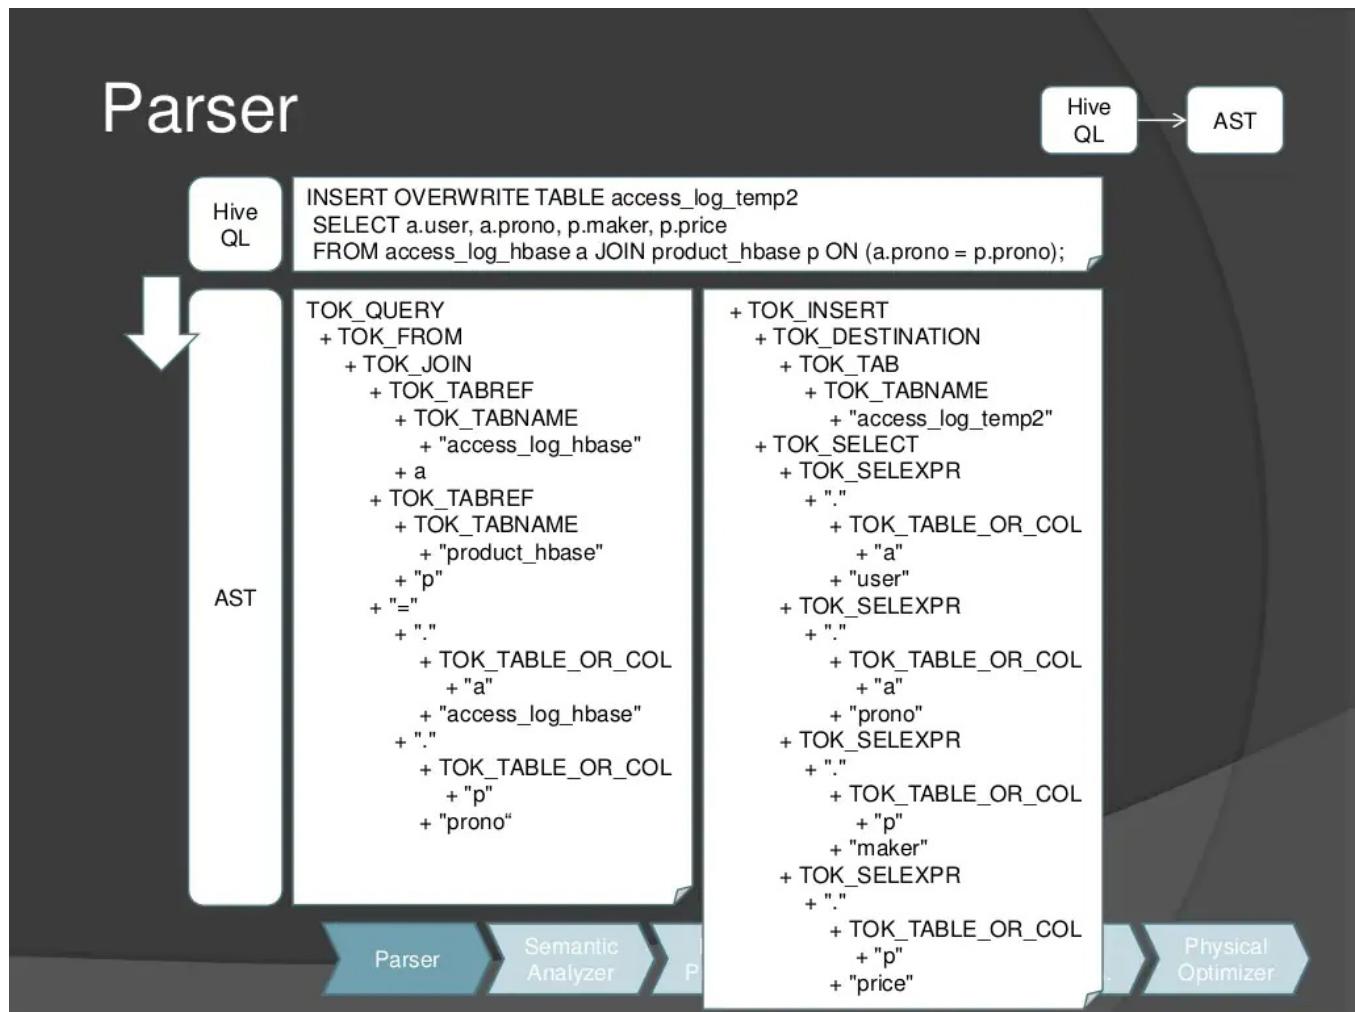


13

# Compiler Overview

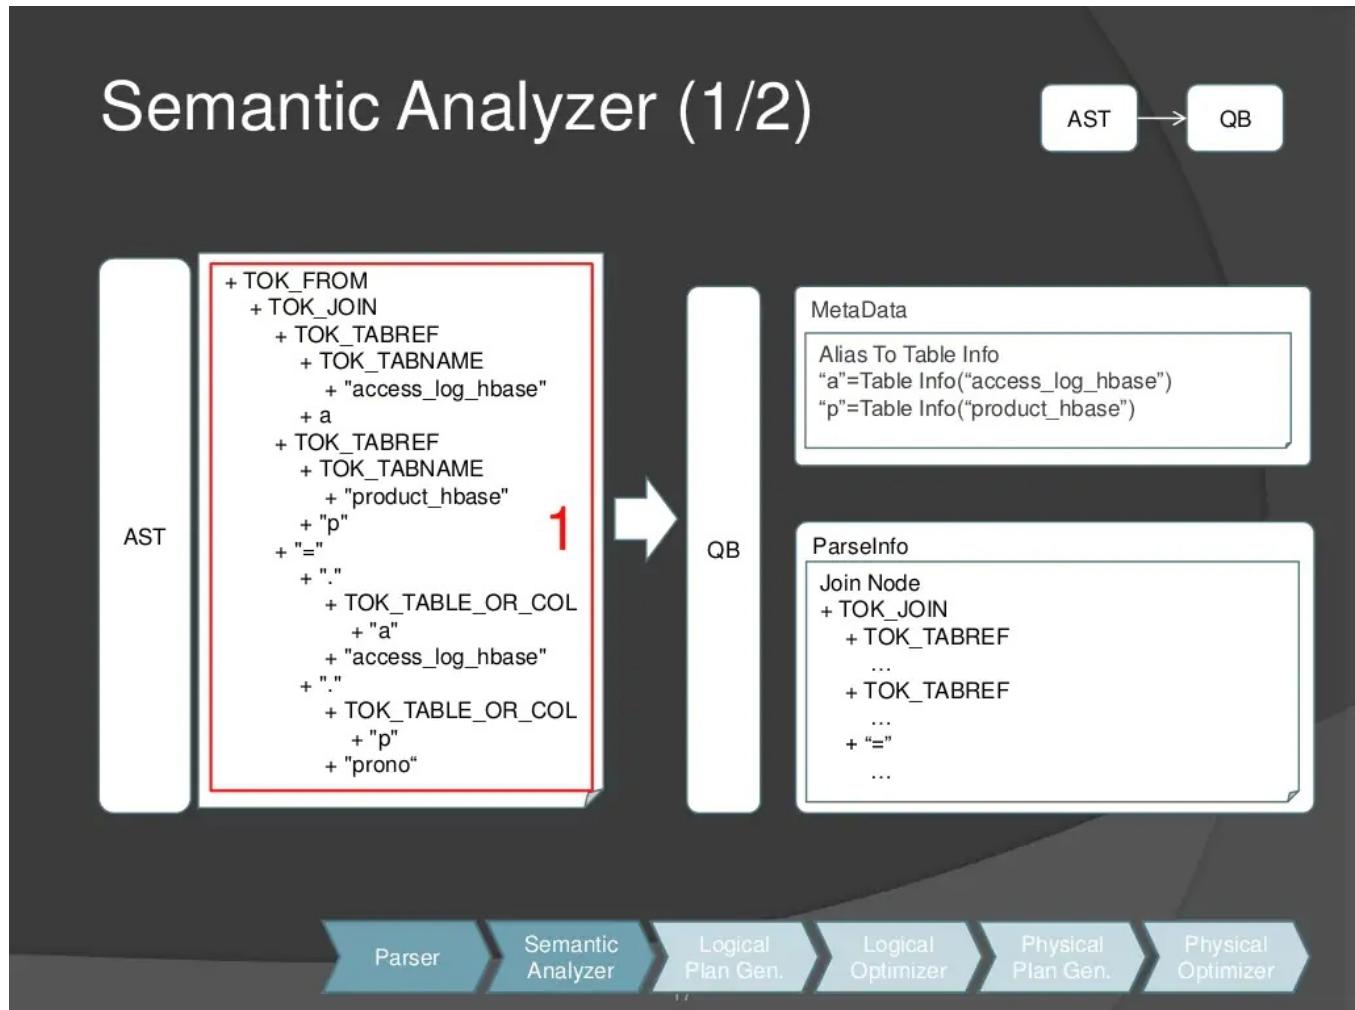


14



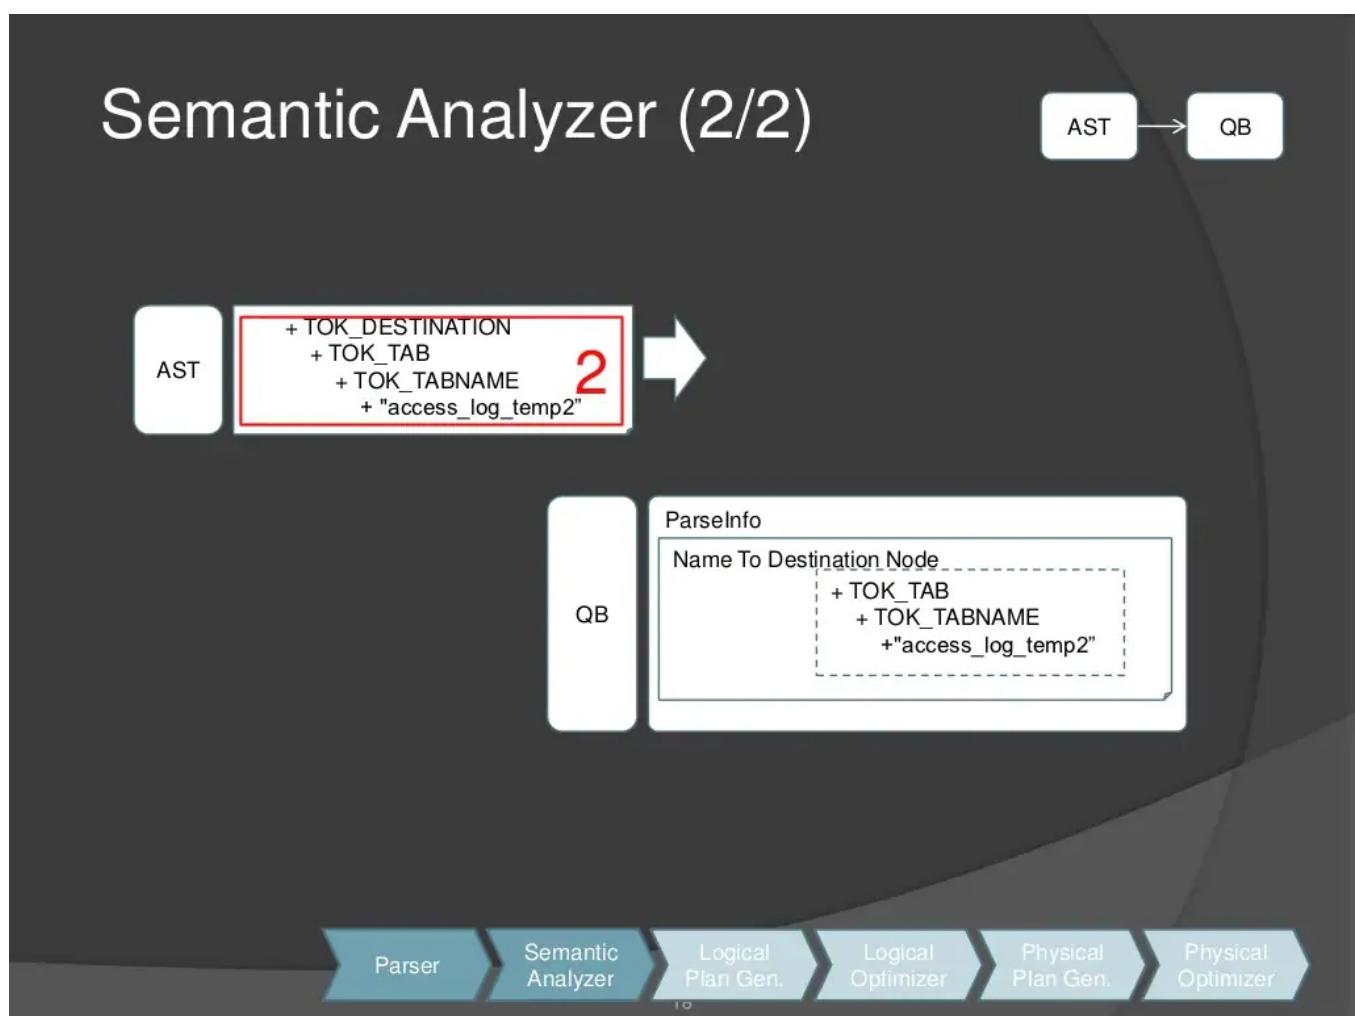
## Semantic Analyzer (1/2)

AST → QB



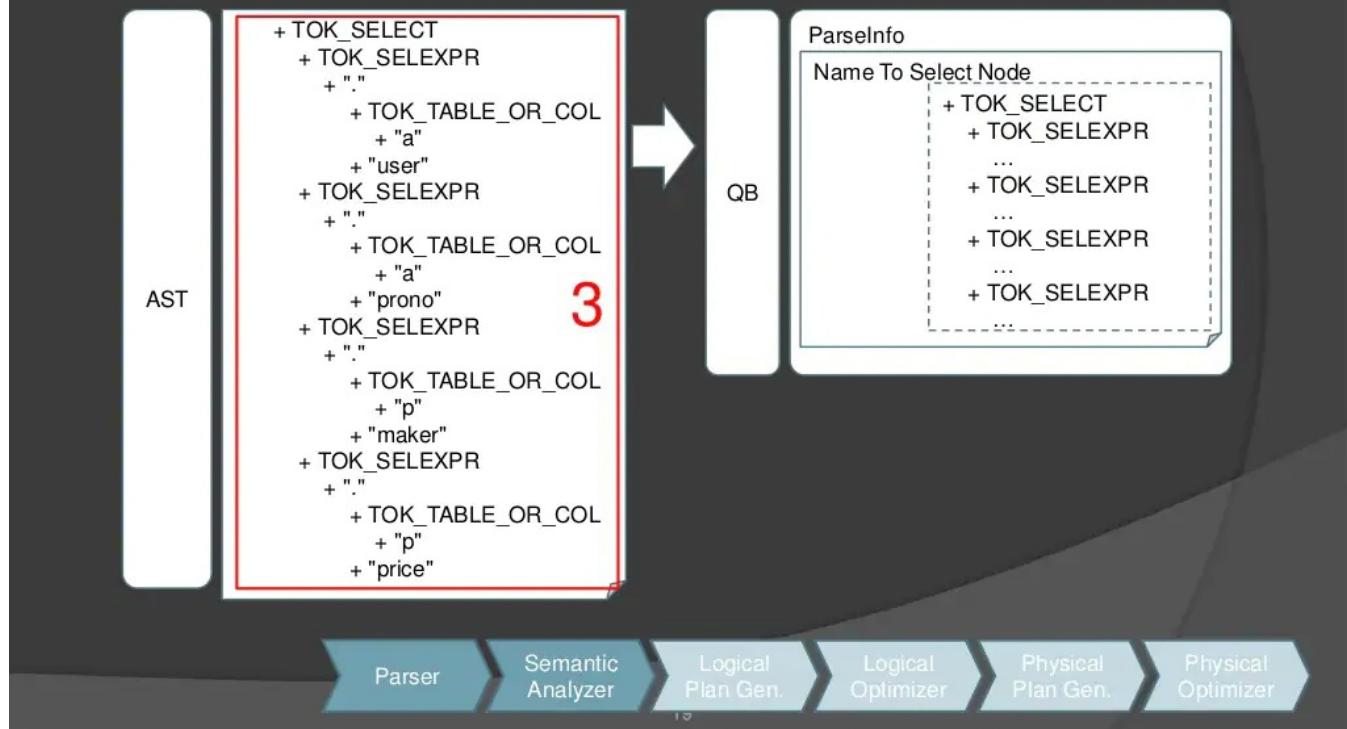
## Semantic Analyzer (2/2)

AST → QB



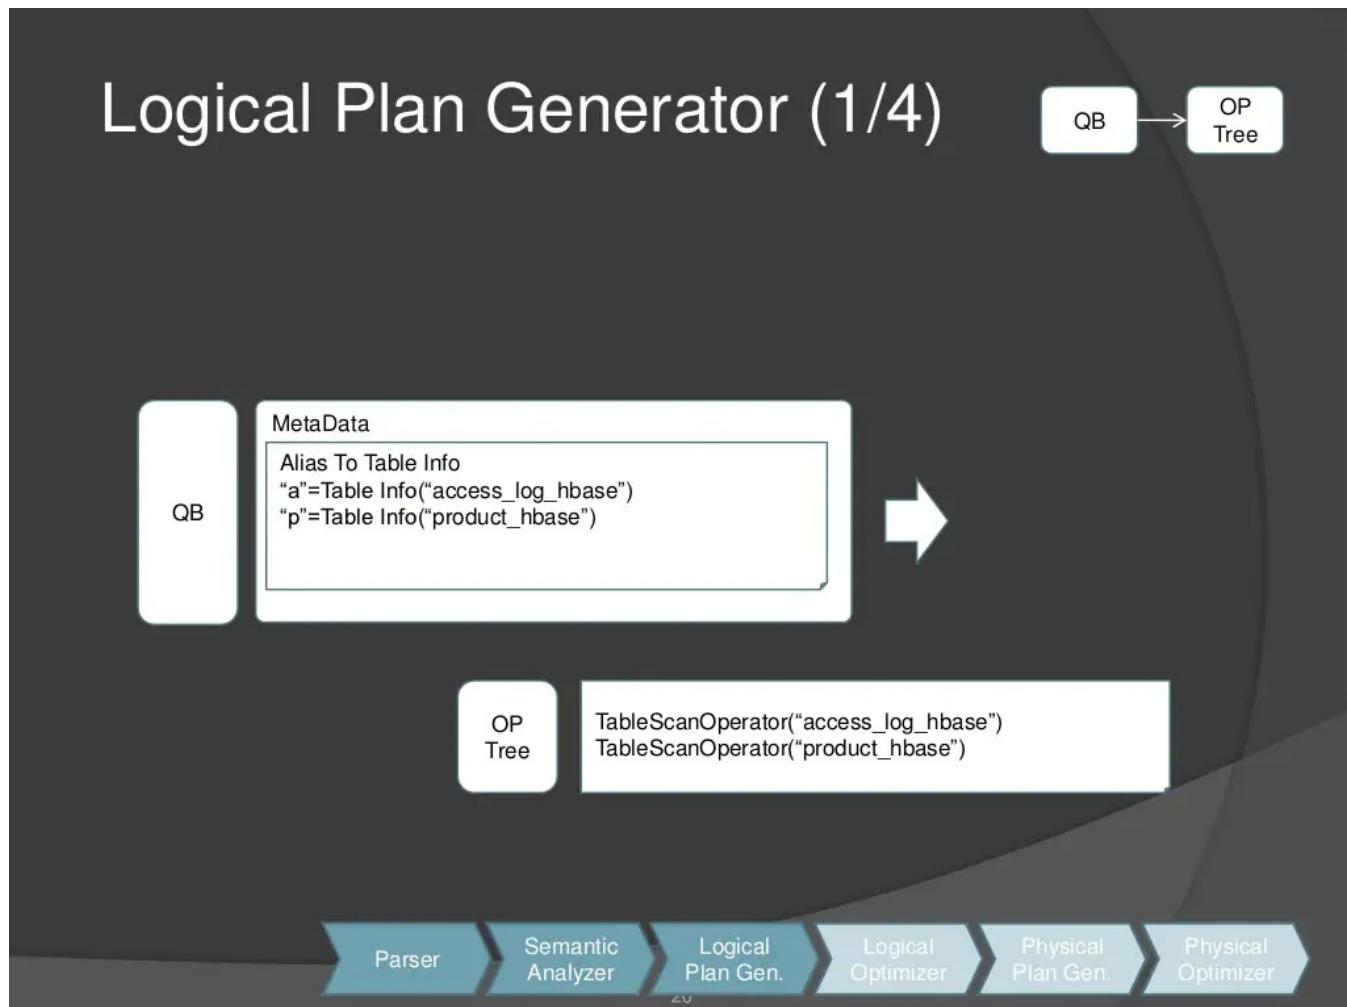
## Semantic Analyzer (2/2)

AST → QB

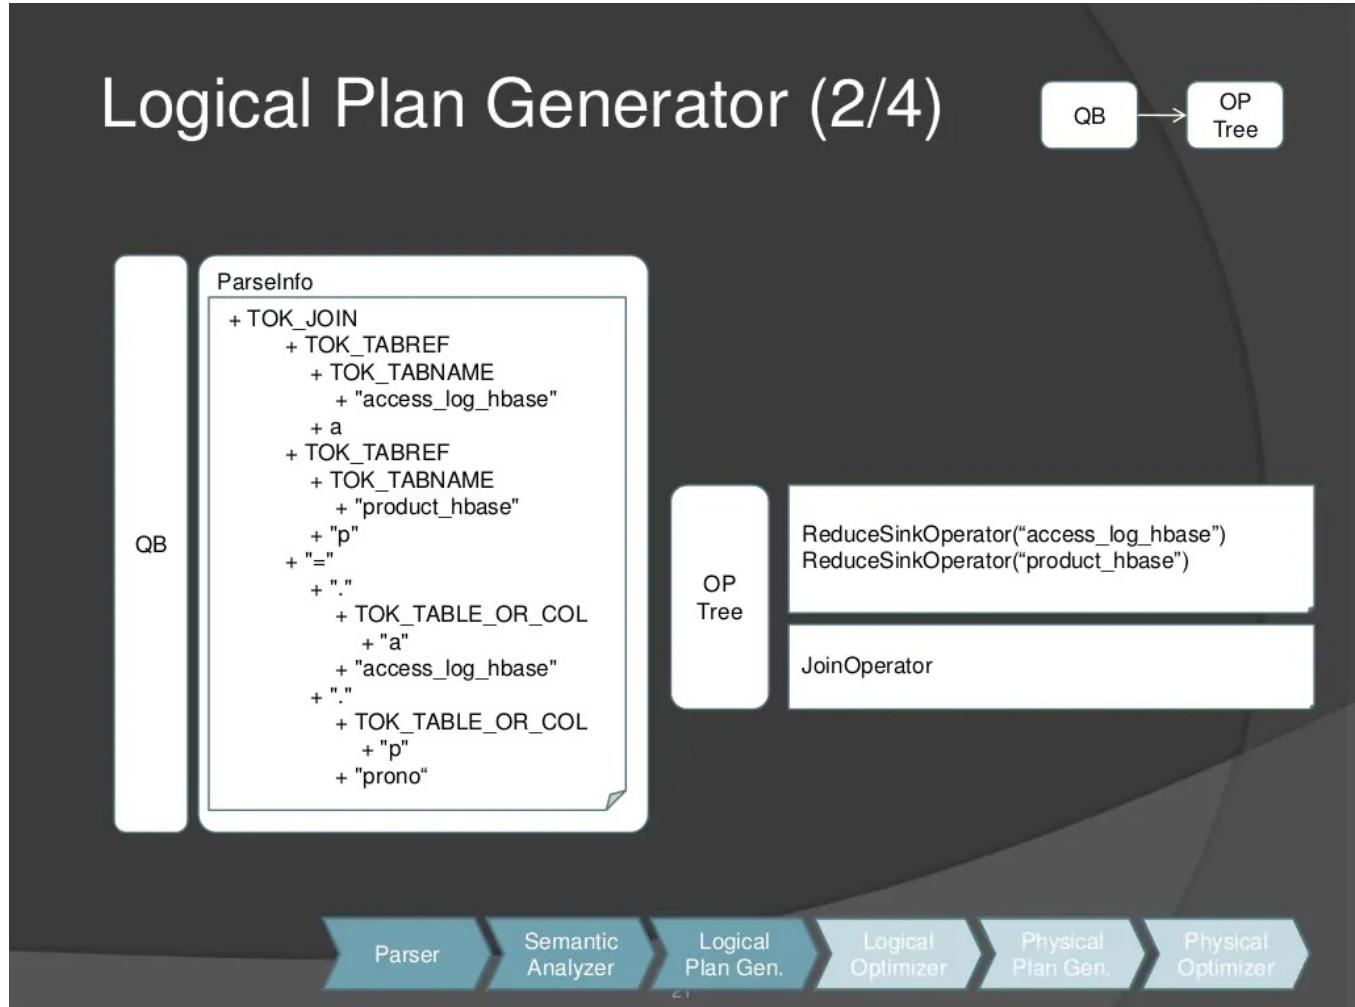


## Logical Plan Generator (1/4)

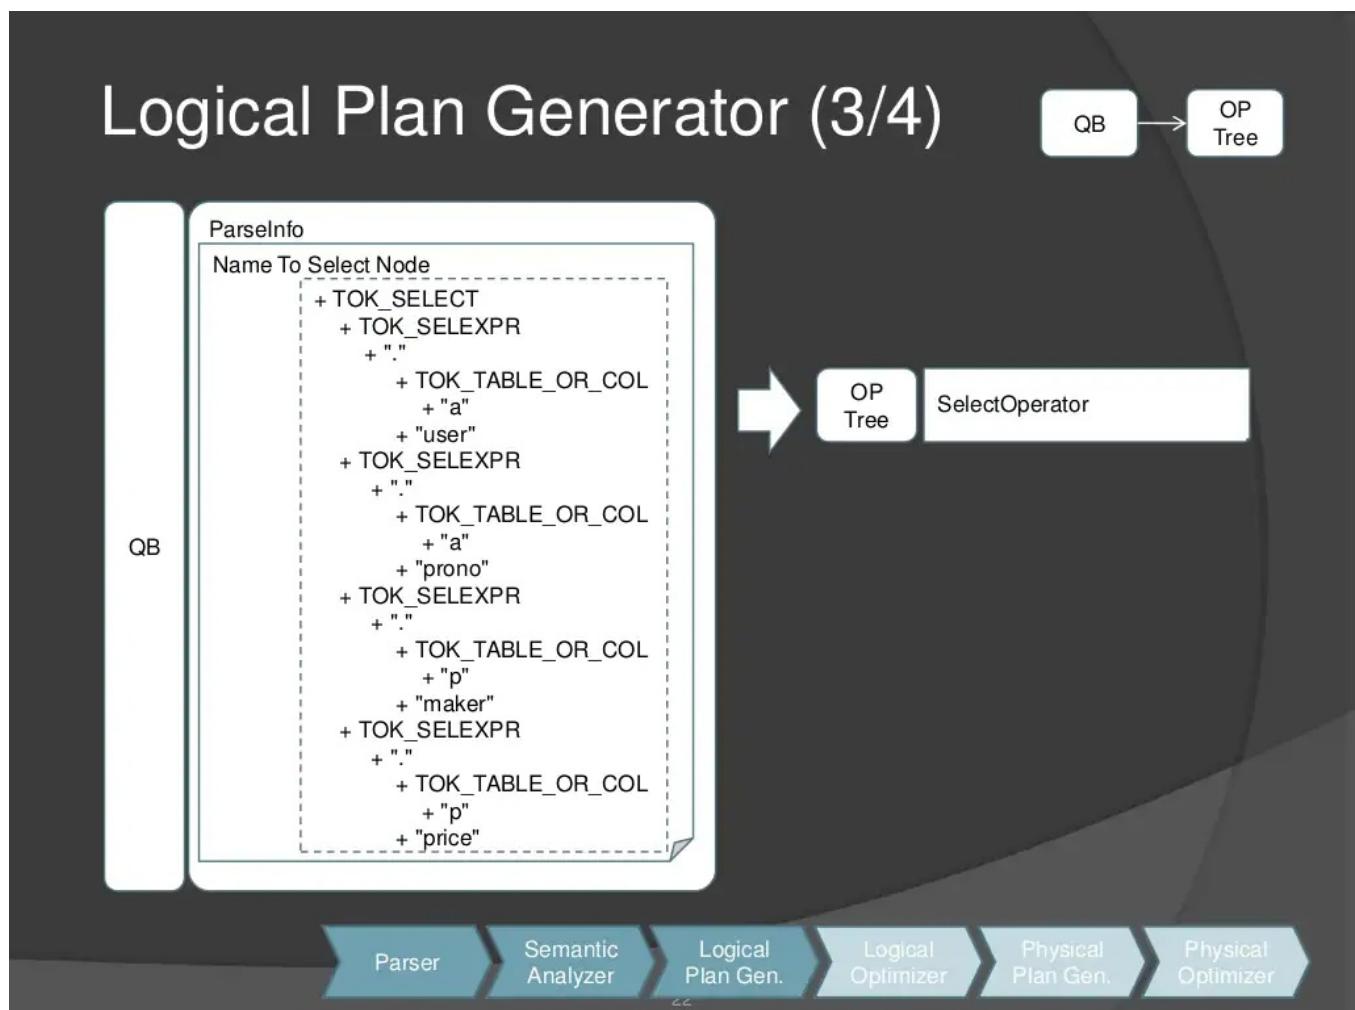
QB → OP Tree



## Logical Plan Generator (2/4)

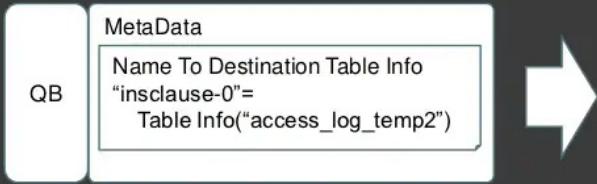


# Logical Plan Generator (3/4)



## Logical Plan Generator (4/4)

QB → OP Tree



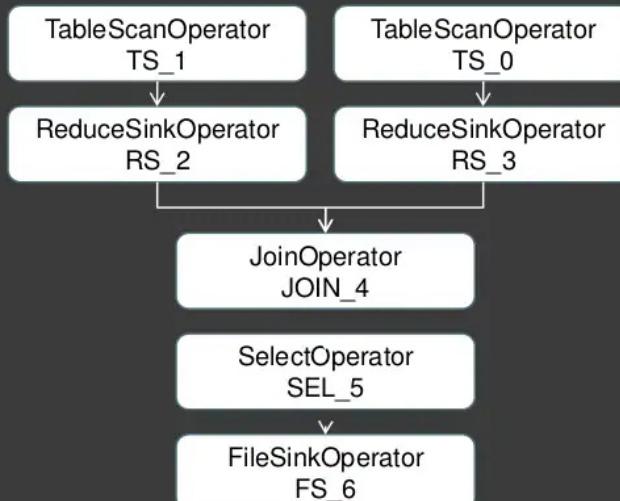
OP Tree

FileSinkOperator

Parser → Semantic Analyzer → Logical Plan Gen. → Logical Optimizer → Physical Plan Gen. → Physical Optimizer

## Logical Plan Generator (result)

OP Tree



Parser → Semantic Analyzer → Logical Plan Gen. → Logical Optimizer → Physical Plan Gen. → Physical Optimizer

# Logical Optimizer

	概要		概要
LineageGenerator	各Operatorが操作するカラムの情報を設定する。	GroupByOptimizer	Group byでの最適化を行う？
ColumnPruner	出力に応じて、各Operatorが入出力するカラムを絞り込む。	SamplePruner	Samplingでの最適化を行う？
Predicate PushDown	条件に応じて、Filterの順番を最適化する。	MapJoinProcessor	MAPJOINが指定されていた場合に、JoinOperatorをMapJoinOperatorにコンバートする。
PartitionPruner	条件に応じて、入力対象となるパーティションを絞り込む。	BucketMapJoin Optimizer	
PartitionCondition Remover	PartitionPrunerで対象のパーティションを絞り込む際に、元となったOperationを削除する。	SortedMergeBucket MapJoinOptimizer	
		UnionProcessor	Identify if both the subqueries of UNION are map-only.
		JoinReader	/*+ STREAMTABLE(A) */
		ReduceSink DeDuplication	If two reducer sink operators share the same partition/sort columns, we should merge them.

Parser

Semantic Analyzer

Logical Plan Gen.

Logical Optimizer

Physical Plan Gen.

Physical Optimizer

25

## Logical Optimizer (Predicate Push Down)

```
INSERT OVERWRITE TABLE access_log_temp2
SELECT a.user, a.prono, p.maker, p.price
FROM access_log_hbase a JOIN product_hbase p ON (a.prono = p.prono);
```



```
INSERT OVERWRITE TABLE access_log_temp2
SELECT a.user, a.prono, p.maker, p.price
FROM access_log_hbase a JOIN product_hbase p ON (a.prono = p.prono)
WHERE p.maker = 'honda';
```

Parser

Semantic Analyzer

Logical Plan Gen.

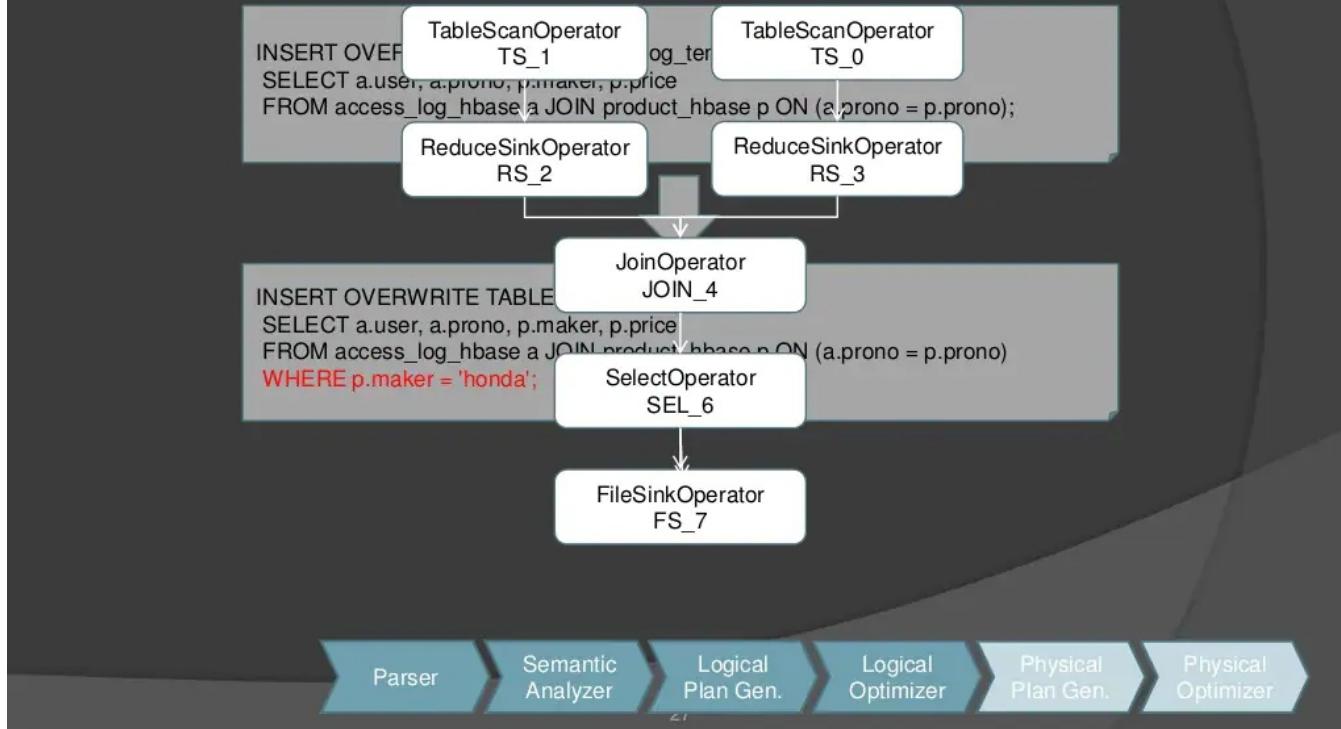
Logical Optimizer

Physical Plan Gen.

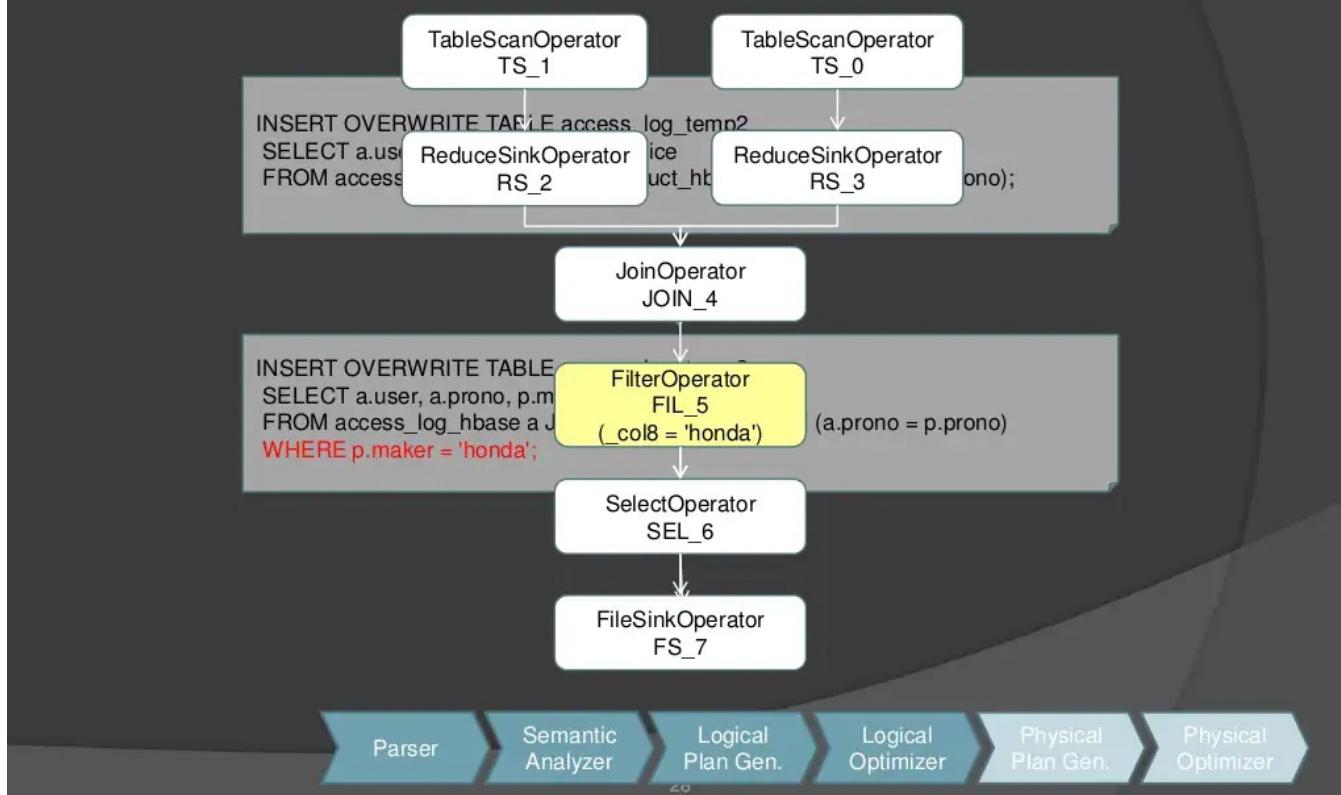
Physical Optimizer

26

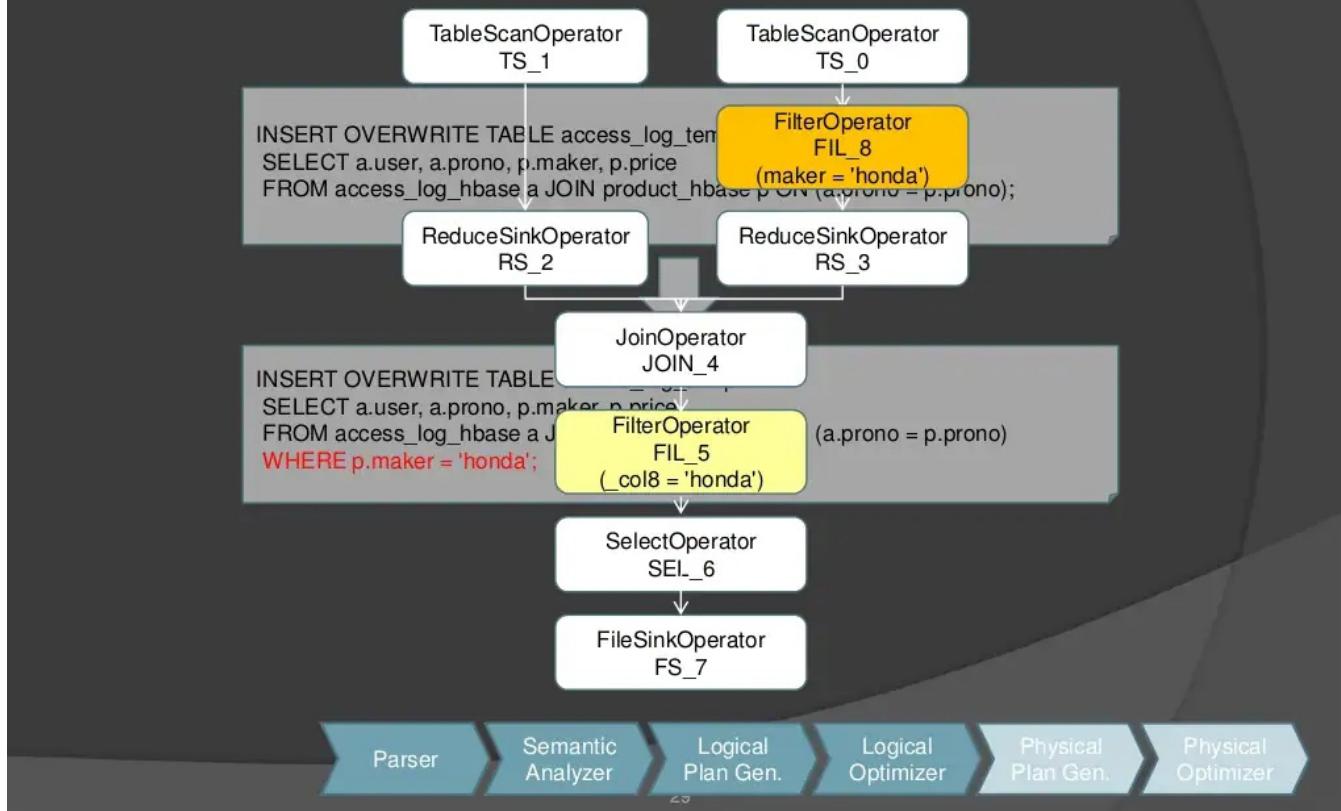
# Logical Optimizer (Predicate Push Down)



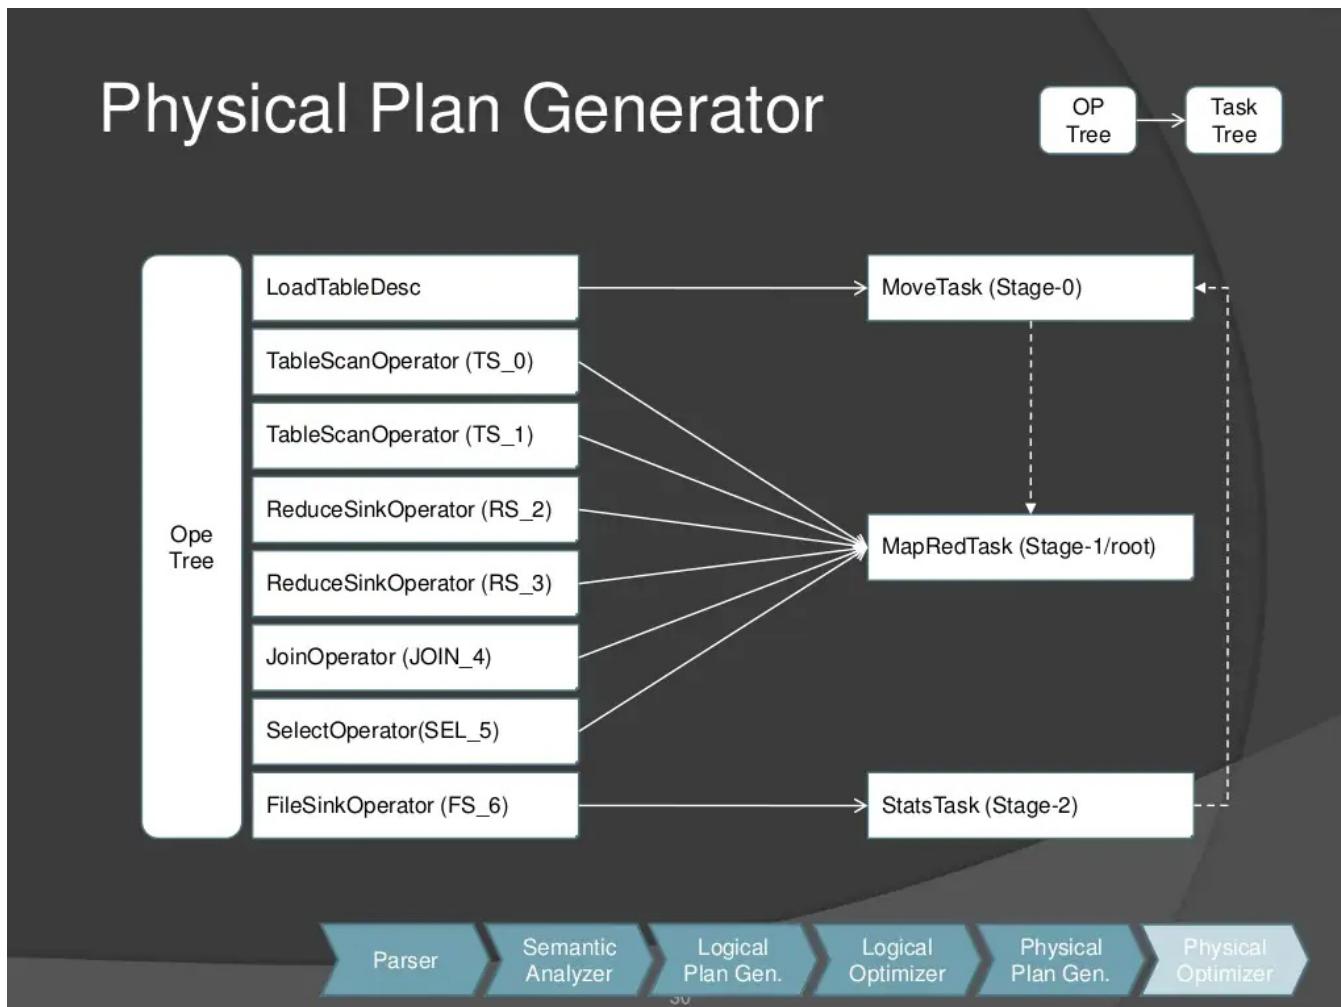
# Logical Optimizer (Predicate Push Down)



# Logical Optimizer (Predicate Push Down)

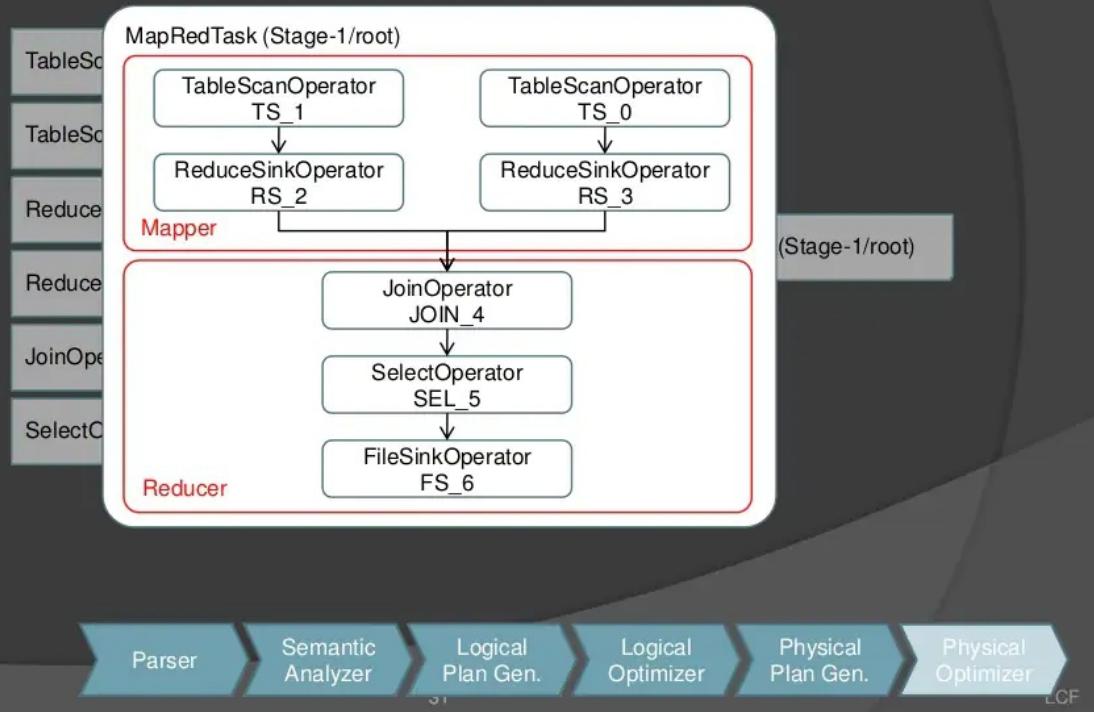


# Physical Plan Generator



# Physical Plan Generator (result)

OP Tree → Task Tree



# Physical Optimizer

Task Tree → Task Tree

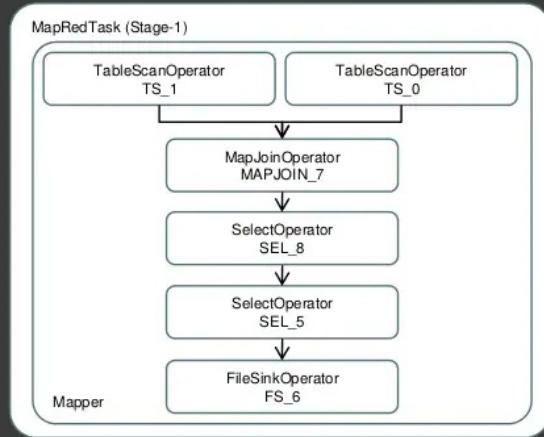
[java/org/apache/hadoop/hive/ql/optimizer/physical/](http://java/org/apache/hadoop/hive/ql/optimizer/physical/)以下

	概要
MapJoinResolver	MapJoinの処理を、ローカルM/Rでのキャッシュ化と実JOINのセット、に変換。
SkewJoinResolver	キーが偏っているテーブルでのJOIN処理を高速化する
CommonJoinResolver	MapJoinへの変換を行う？

Parser → Semantic Analyzer → Logical Plan Gen. → Logical Optimizer → Physical Plan Gen. → Physical Optimizer

# Physical Optimizer (MapJoinResolver)

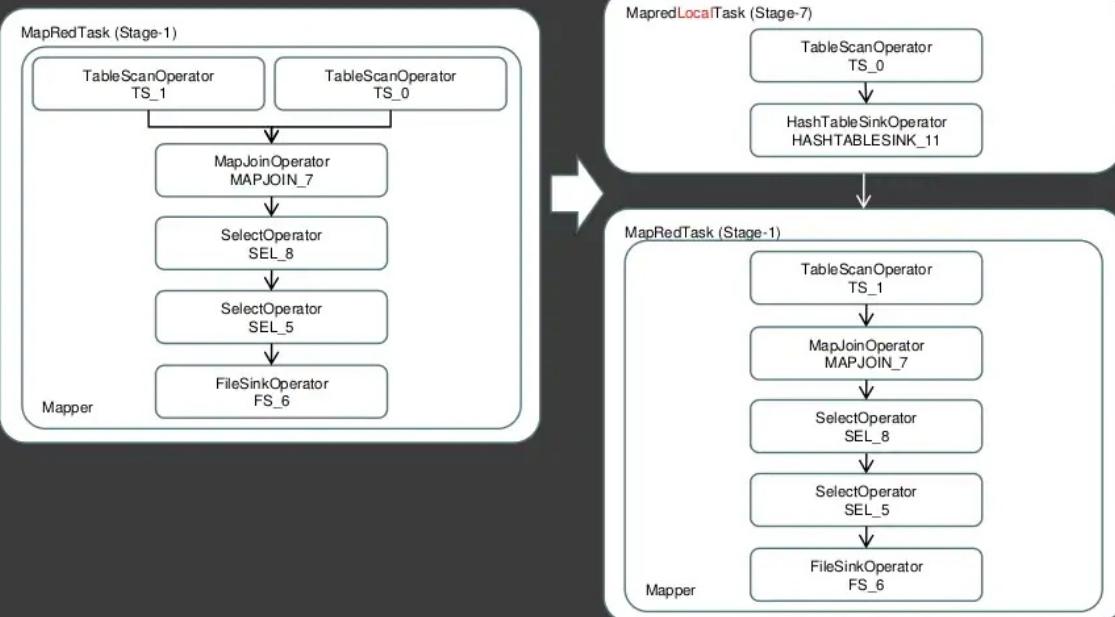
Task Tree → Task Tree



Parser → Semantic Analyzer → Logical Plan Gen. → Logical Optimizer → Physical Plan Gen. → Physical Optimizer

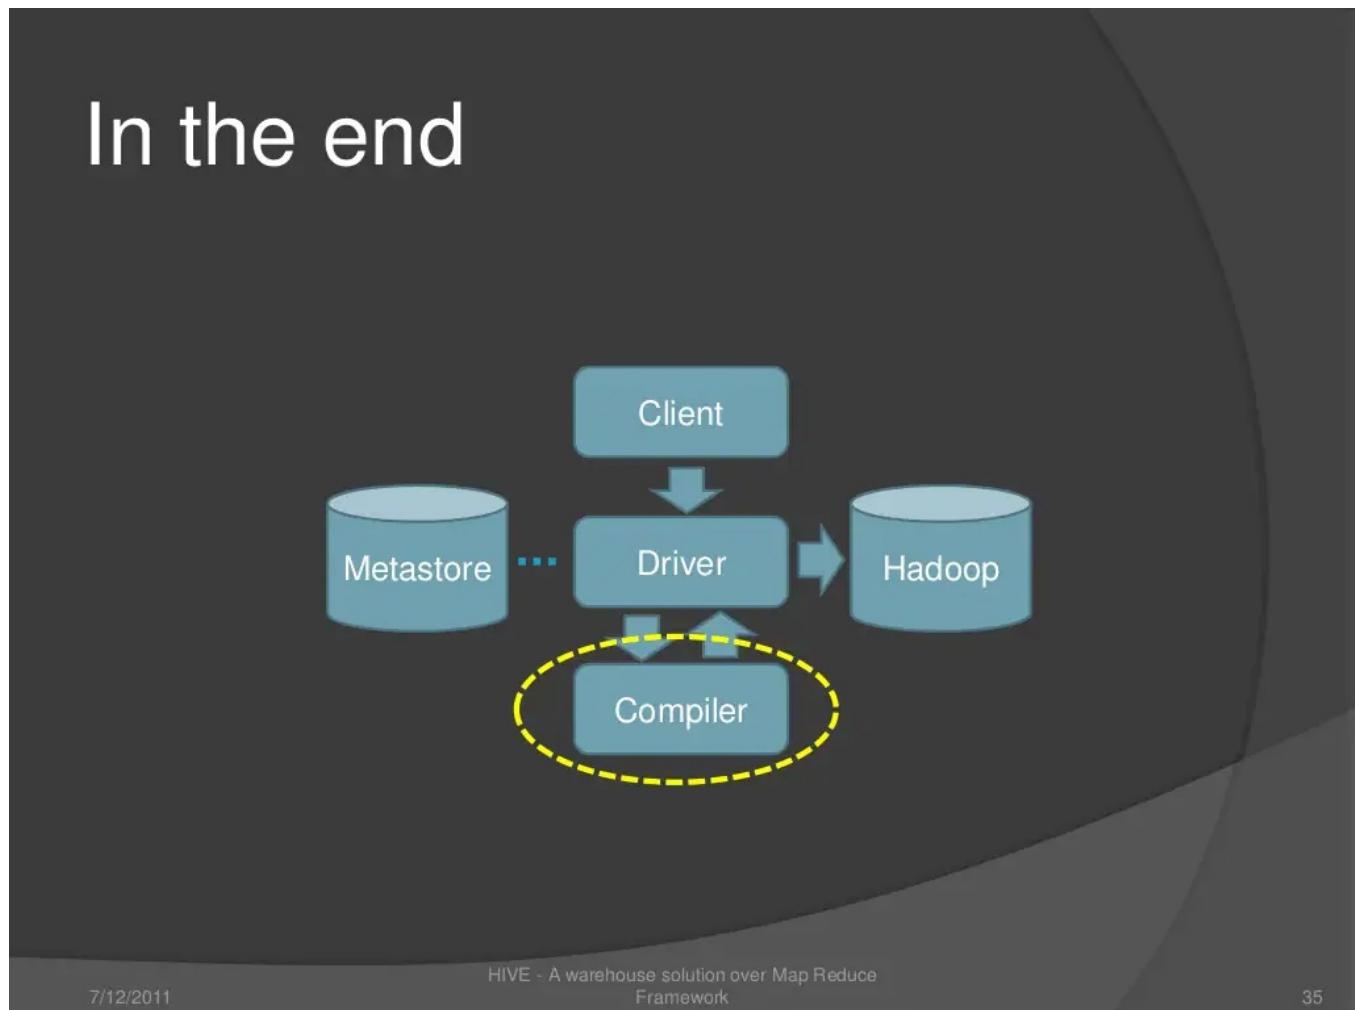
# Physical Optimizer (MapJoinResolver)

Task Tree → Task Tree

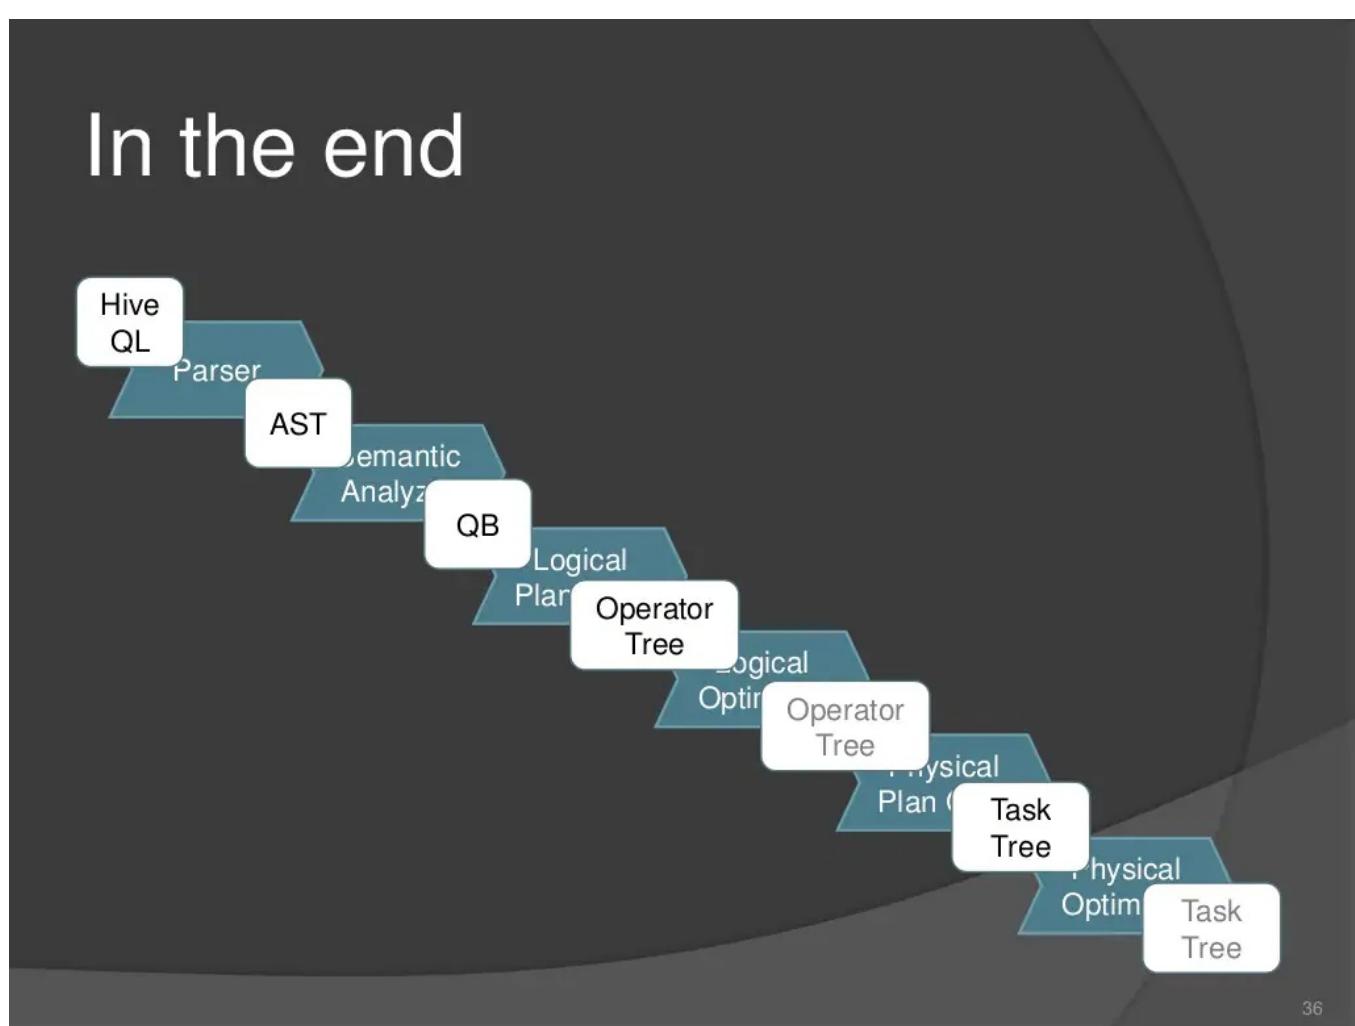


Parser → Semantic Analyzer → Logical Plan Gen. → Logical Optimizer → Physical Plan Gen. → Physical Optimizer

# In the end



# In the end



End

- ➊ Appendix: What does Explain show?



# Appendix: What does Explain show?

# Appendix: What does Explain show?

# Appendix: What does Explain show?

